YANG Runming

(+86) 17362995487 ◇ yang3121099@126.com ◇ https://rummyyang.github.io/

Research interests: LLM post-training; Knowledge distillation; Efficient reasoning; Low-bit quantization

EDUCATION

The University of Hong Kong Oct. 2025 - Present
Ph.D. student in Electrical and Electronic Engineering — Advisor: Prof. Ngai Wong; expected Sep. 2029

Tsinghua University Sep. 2022 - Jun. 2025
M.Eng. in Big Data Engineering — Advisor: Prof. Yujiu Yang; GPA: 3.73/4.00

Huazhong University of Science and Technology Sep. 2017 - Jun. 2021
B.Eng. in AI and Automation — Science and Engineering Cross-Innovation Experimental Class

PUBLICATIONS

* Equal contribution.

[1] T. Wu*, **R. Yang***, T. Liu, J. Wang, N. Wong. *Revisiting Model Interpolation for Efficient Reasoning.*
ACL 2026 (Main) [Paper] [Code]

[2] T. Liu*, T. Wu*, **R. Yang**, S. Sun, J. Wang, Y. Yang. *ProFit: Leveraging High-Value Signals in SFT via Probability-Guided Token Selection.*
ACL 2026 (Findings) [Paper] [Code]

[3] T. Wu*, Y. Cheng*, C. Ding*, **R. Yang**, X. Feng, W. Zhou, Z. Liu, N. Wong. *Can We Trust LLMs on Memristors? Diving into Reasoning Ability under Non-Ideality.*
COLM 2026 [Paper]

[4] H. Xiao, Q. Yang, D. Xie, W. Xu, Z. Su, **R. Yang**, H. Liu, W. Zhou, Z. Liu, N. Wong. *Exploring Layer-wise Information Effectiveness for Post-Training Quantization in Small Language Models.*
ACL 2026 (Findings) [Paper] [Project]

[5] T. Wu, C. Tao, J. Wang, **R. Yang**, Z. Zhao, N. Wong. *Rethinking Kullback-Leibler Divergence in Knowledge Distillation for Large Language Models.*
COLING 2025 (Oral) [Paper] [Code]

[6] **R. Yang**, T. Wu, Y. Yang. *LoCa: Logit Calibration for Knowledge Distillation.*
ECAI 2024 [Paper]

RESEARCH INTERNSHIP

Tencent AI Lab Nov. 2023 - Jun. 2025
Research Intern, Information Security Department

- Synthesized multimodal data and conducted model post-training and evaluation for image-text content moderation in QQ group chats, Qzone, and short-video comments.
- Worked on tiered deployment and inference acceleration; trained models were deployed in internal risk-control systems, improving throughput and reducing false-positive rates over the previous solution.
- Researched large/small model collaboration, data distillation, and lightweight modeling; contributed to two academic papers and one company patent.

ONGOING RESEARCH

CoT Data and Efficient Reasoning Post-training Ongoing

Collaborative Ph.D. research at the University of Hong Kong

- Planned approach: compare offline CoT fine-tuning with on-policy distillation on student-generated trajectories, examining data quality and teacher-student distribution alignment.
- Research objective: improve reasoning quality and efficiency under a controlled compute budget, with planned evaluation of accuracy, reasoning length, and training cost; manuscript in preparation.

PREPRINTS AND SUBMITTED MANUSCRIPTS

* Equal contribution.

[P1] **R. Yang***, T. Wu*, J. Li, P. Hu, Y.-C. Wu, N. Wong, Y. Yang. *Shadow-FT: Tuning Instruct Model via Training on Paired Base Model.*
Submitted to NAACL 2027 [Paper] [Code]

[P2] **R. Yang***, T. Wu*, J. Wang, P. Hu, Y.-C. Wu, N. Wong, Y. Yang. *LLM-NEO: Parameter Efficient Knowledge Distillation for Large Language Models.*
arXiv:2411.06839 [Paper] [Code]

[P3] Q. Yang, **R. Yang**, H. Xiao, W. Xu, J. Chen, H. Liu, C. Ding, R. Hu, Y.-C. Wu, N. Wong. *FluxBin: Flexible LUT-based Ultra-low-bit LLM Inference by Algorithm-Kernel Synergy.*
Submitted to NAACL 2027 [Paper] [Code]

[P4] T. Wu, **R. Yang**, T. Liu, J. Wang, Z. Xu, N. Wong. *Timber: Training-free Instruct Model Refining with Base Model via Effective Rank.*
Submitted to NAACL 2027 [Paper] [Code]

[P5] H. Xiao*, **R. Yang***, Q. Yang*, W. Xu, Z. Li, Y. Su, Z. Liu, H. Yang, N. Wong. *PTQTP: Post-Training Quantization to Trit-Planes for Large Language Models.*
arXiv:2509.16989 [Paper] [Project]

[P6] Y. Zhang, T. Wu, **R. Yang**, T. Liu, J. Wang, Y. Yang. *From Answers to Obligations: Benchmarking Financial Instruction Following with FinIF.*
Submitted to EMNLP 2027

[P7] Z. Li*, Y. Su*, S. Wang, **R. Yang**, C. Xie, A. Liu, M. Li, J. Cao, Y. Xie, N. Wong, H. Yang. *Quantization Meets Reasoning: Exploring and Mitigating Degradation of Low-Bit LLMs in Mathematical Reasoning.*
arXiv:2505.11574 [Paper]

[P8] T. Liu*, T. Wu*, M. Zheng, X. Luo, **R. Yang**, X. Yang, J. Wang, Y. Yang. *Simple-OPD: Demystifying Warm-up for On-policy Distillation.*
Submitted to NAACL 2027 [Paper] [Code]